

Neural Networks

Steve Cygu

January 24, 2019

OUTLINE

Introduction

- Why Machine Learning?
- Neural Networks and Human Brain

Neural Networks

- Types of Neural Networks
- Components

Fitting Neural Networks

- Feed-forward
- Back propagation
- Gradient descent

Some Issues in Training Neural Networks

Example

- Methods
- results

References


Why Machine Learning?

Can we write algorithm to correctly identify each of the objects?


504192



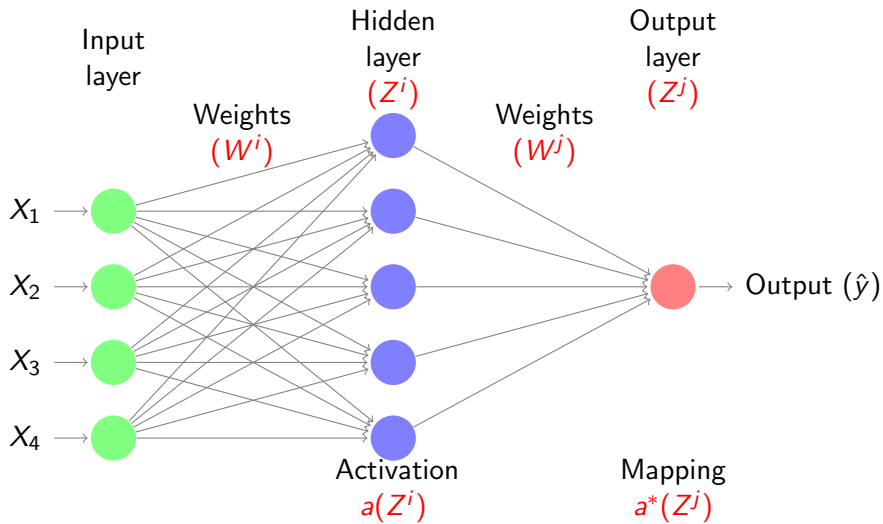
Neural Networks and Human Brain



The Neuron and information transmission




Multi-layer perceptron (MLP)




Single-layer Perceptron

- ▶ No hidden layer, a single neuron.



Single-layer Perceptron

- No hidden layer, a single neuron.



$$\hat{y} = b + \sum_{i=1}^n x_i w_i$$

Components

Layers

- ▶ **Input nodes**
No computation
- ▶ **Hidden nodes (Neurons)**
Intermediate processing, computation and transfers to another hidden layer or output.
- ▶ **Output nodes**
Uses a function (not necessarily activation function) to map the input from other layers to desired output format.
 - ▶ Sigmoid
 - ▶ Softmax

Synapse/Connections

- ▶ Transfers the output of neuron i to the input of neuron j .
- ▶ Each connection is assigned weight, W_{ij}

Activation function

Introduces **nonlinearity** into the neuron output.

- ▶ Sigmoid (Logistic Activation Function)

$$a(z) = \frac{1}{1 + \exp(-z)}$$



- ▶ Tanh (hyperbolic tangent Activation Function)

$$a(z) = \tanh(z) = \frac{2}{1 + \exp(-2z)} - 1 = 2\text{sigmoid}(2z) - 1$$

- ▶ ReLU (Rectified Linear Unit Activation Function)

$$a(z) = \max(0, z)$$

Sigmoid and Tanh Activation Functions



Fitting Neural Networks

Weights are the parameters. The generic approach is by **gradient descent**.

- ▶ Forward-propagation (feed-forward)
- ▶ Backward-propagation

Feed-forward

- ▶ Randomly assign starting weights and consider any Z_k^i , $k = 1 \cdots n$. Then,

Feed-forward

- ▶ Randomly assign starting weights and consider any Z_k^i , $k = 1 \cdots n$. Then,

$$Z_1^i = X_1 W_{1,1}^i + X_2 W_{2,1}^i + X_3 W_{3,1}^i + X_4 W_{4,1}^i$$

Feed-forward

- ▶ Randomly assign starting weights and consider any Z_k^i , $k = 1 \cdots n$. Then,

$$Z_1^i = X_1 W_{1,1}^i + X_2 W_{2,1}^i + X_3 W_{3,1}^i + X_4 W_{4,1}^i$$

$$Z_2^i = X_1 W_{1,2}^i + X_2 W_{2,2}^i + X_3 W_{3,2}^i + X_4 W_{4,2}^i$$

Feed-forward

- ▶ Randomly assign starting weights and consider any Z_k^i , $k = 1 \cdots n$. Then,

$$Z_1^i = X_1 W_{1,1}^i + X_2 W_{2,1}^i + X_3 W_{3,1}^i + X_4 W_{4,1}^i$$

$$Z_2^i = X_1 W_{1,2}^i + X_2 W_{2,2}^i + X_3 W_{3,2}^i + X_4 W_{4,2}^i$$

\vdots

Feed-forward

- ▶ Randomly assign starting weights and consider any Z_k^i , $k = 1 \cdots n$. Then,

$$Z_1^i = X_1 W_{1,1}^i + X_2 W_{2,1}^i + X_3 W_{3,1}^i + X_4 W_{4,1}^i$$

$$Z_2^i = X_1 W_{1,2}^i + X_2 W_{2,2}^i + X_3 W_{3,2}^i + X_4 W_{4,2}^i$$

\vdots

Z^i 'component' is the **sum of weighted inputs** to each neuron.

$$Z^i = XW^i \tag{1}$$

Feed-forward

Apply activation function to 1

$$a^i = a(Z^i) \quad (2)$$

Propagate 2 to the output layer

$$Z^j = a^i W^j \quad (3)$$

$$\implies \hat{y} = a^j = a^*(Z^j) \quad (4)$$

Back propagation

- ▶ Aim is to estimate weights that ensures the model fits the training data well.
- ▶ Calculate the error at the output nodes and propagate them back to the network.

$$J = \sum \frac{1}{2}(y - \hat{y})^2 \quad (5)$$

Back propagation

- ▶ Aim is to estimate weights that ensures the model fits the training data well.
- ▶ Calculate the error at the output nodes and propagate them back to the network.

$$J = \sum \frac{1}{2} (y - \hat{y})^2 \quad (5)$$

$$J(W) = \frac{1}{2} \sum (y - a^*(a(XW^i)W^j))^2 \quad (6)$$

Back propagation

- ▶ Aim is to estimate weights that ensures the model fits the training data well.
- ▶ Calculate the error at the output nodes and propagate them back to the network.

$$J = \sum \frac{1}{2} (y - \hat{y})^2 \quad (5)$$

$$J(W) = \frac{1}{2} \sum (y - a^*(a(XW^i)W^j))^2 \quad (6)$$

- ▶ Compute the gradient; $\frac{\partial J}{\partial W^i}$ and $\frac{\partial J}{\partial W^j}$

Back propagation

- ▶ Aim is to estimate weights that ensures the model fits the training data well.
- ▶ Calculate the error at the output nodes and propagate them back to the network.

$$J = \sum \frac{1}{2} (y - \hat{y})^2 \quad (5)$$

$$J(W) = \frac{1}{2} \sum (y - a^*(a(XW^i)W^j))^2 \quad (6)$$


- ▶ Compute the gradient; $\frac{\partial J}{\partial W^i}$ and $\frac{\partial J}{\partial W^j}$
- ▶ Adjust the weights using optimization method such as **Gradient Descent**.

Gradient descent

$$W_{t+1} = W_t - \gamma \Delta J(W_t)$$

Gradient descent

$$W_{t+1} = W_t - \gamma \Delta J(W_t)$$



Some Issues in Training Neural Networks

1. Starting values

- ▶ Starting weights are random numbers near zero.
- ▶ However, near zero weights collapses NN into approximately linear model.
- ▶ Exactly zero weights leads to zero derivatives and perfect symmetry.
- ▶ Large weights lead to poor results.

Some Issues in Training Neural Networks


1. Starting values

- ▶ Starting weights are random numbers near zero.
- ▶ However, near zero weights collapses NN into approximately linear model.
- ▶ Exactly zero weights leads to zero derivatives and perfect symmetry.
- ▶ Large weights lead to poor results.

2. Overfitting and Stopping Criterion

- ▶ Reduce training error to some predetermined threshold → overfitting.
- ▶ Regularization by *weight decay* (analogous to ridge regression), $\lambda \geq 0$. Larger values of λ shrinks weights toward zero.
- ▶ Cross-validation is used to estimate λ .

3. Convergence at the Local Minima



Example

- ▶ Breast Cancer Wisconsin Data Set

Example

- ▶ Breast Cancer Wisconsin Data Set
- ▶ The data set contains 569 cases with 31 variables. The diagnosis classification is either (M = Malignant) or (B = Benign).

Example

- ▶ Breast Cancer Wisconsin Data Set
- ▶ The data set contains 569 cases with 31 variables. The diagnosis classification is either (M = Malignant) or (B = Benign).

Show 32 entries

Search:

	vars	labels
1	id_number	ID number
2	diagnosis	Diagnosis (M = malignant, B = benign)
3	radius_mean	radius (mean of distances from center to points on the perimeter) - Mean
4	texture_mean	texture (standard deviation of gray-scale values) - Mean
5	perimeter_mean	perimeter - Mean
6	area_mean	area - Mean
7	smoothness_mean	smoothness (local variation in radius lengths) - Mean
8	compactness_mean	compactness (perimeter ² / area - 1.0) - Mean
9	concavity_mean	concavity (severity of concave portions of the contour) - Mean
10	concave_points_mean	concave points (number of concave portions of the contour) - Mean
11	symmetry_mean	symmetry - Mean
12	fractal_dimension_mean	fractal dimension ("coastline approximation" - 1) - Mean
13	radius_se	radius (mean of distances from center to points on the perimeter) - Standard error
14	texture_se	texture (standard deviation of gray-scale values) - Standard error
15	perimeter_se	perimeter - Standard error
16	area_se	area - Standard error
17	smoothness_se	smoothness (local variation in radius lengths) - Standard error
18	compactness_se	compactness (perimeter ² / area - 1.0) - Standard error
19	concavity_se	concavity (severity of concave portions of the contour) - Standard error
20	concave_points_se	concave points (number of concave portions of the contour) - Standard error
21	symmetry_se	symmetry - Standard error
22	fractal_dimension_se	fractal dimension ("coastline approximation" - 1) - Standard error
23	radius_worst	radius (mean of distances from center to points on the perimeter) - Worst
24	texture_worst	texture (standard deviation of gray-scale values) - Worst
25	perimeter_worst	perimeter - Worst
26	area_worst	area - Worst
27	smoothness_worst	smoothness (local variation in radius lengths) - Worst
28	compactness_worst	compactness (perimeter ² / area - 1.0) - Worst
29	concavity_worst	concavity (severity of concave portions of the contour) - Worst
30	concave_points_worst	concave points (number of concave portions of the contour) - Worst

Methods

- ▶ Data partitioning: 80% training set and 20% test set
- ▶ Automatic grid search, with a *tunelength* = 10 was used to find optimal parameter values

Methods

- ▶ Data partitioning: 80% training set and 20% test set
- ▶ Automatic grid search, with a *tunelength* = 10 was used to find optimal parameter values
 - ▶ Weight decay


Methods

- ▶ Data partitioning: 80% training set and 20% test set
- ▶ Automatic grid search, with a *tunelength* = 10 was used to find optimal parameter values
 - ▶ Weight decay
 - ▶ Number of hidden neurons


Methods

- ▶ Data partitioning: 80% training set and 20% test set
- ▶ Automatic grid search, with a *tunelength* = 10 was used to find optimal parameter values
 - ▶ Weight decay
 - ▶ Number of hidden neurons
- ▶ 10-fold cross validation
- ▶ ROC was used as the performance metric

results : Weight decay and hidden neurons




results : Weight decay and hidden neurons

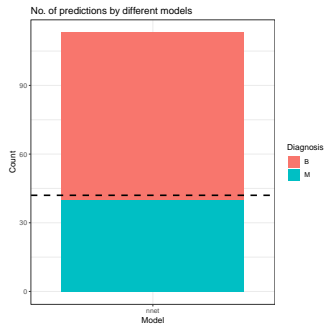


	size	decay
1	3	0.04216965


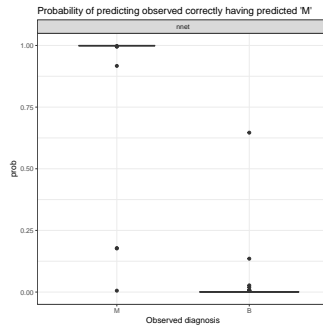
Result: Resampling distribution



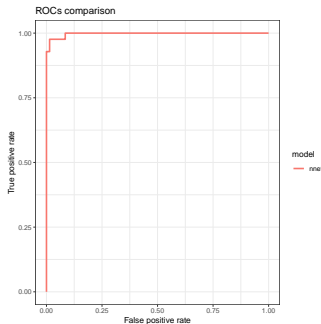

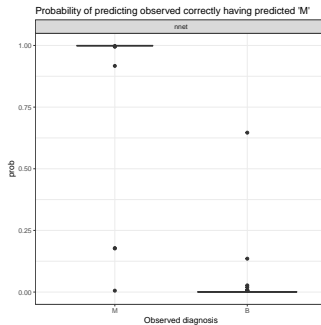
Result: Predictions



Result: Predictions



Result: Predictions



Conclusions

- ▶ Discussed Gradient descent in neural network
- ▶ Applied neural network to classify breast cancer. The model provides a good classification of the data, with ROC of 0.99.

References

- [1] Trevor, H., Robert, T., & JH, F. (2009). The elements of statistical learning: data mining, inference, and prediction. *Springer series in statistics*. Second Edition
- [2] Tom M. Mitchell. (1997). Machine Learning *McGraw-Hill International Editions*.
- [3] Internet sources (2019).